

```

# Project: Fine-Grained Bird Species Classification on CUB-200-2011
# Compare: Pretrained ResNet-18 vs Pretrained ViT

# =====
# 1. Install dependencies (if needed in Colab)
# =====
!pip install -q transformers datasets torch torchvision scikit-learn matplotlib pillow tqdm evaluate
!pip install seaborn
# =====
# 2. Import libraries
# =====

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, random_split
from torchvision import models, transforms
from transformers import ViTForImageClassification
from datasets import load_dataset
from tqdm import tqdm
from sklearn.metrics import accuracy_score, top_k_accuracy_score, classification_report, confusion_matrix
from collections import Counter
import numpy as np
import matplotlib.pyplot as plt
import time
import seaborn as sns
import pandas as pd

Requirement already satisfied: seaborn in /usr/local/lib/python3.12/dist-packages (0.13.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in /usr/local/lib/python3.12/dist-packages (from seaborn) (2.0.2)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.12/dist-packages (from seaborn) (2.2.2)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /usr/local/lib/python3.12/dist-packages (from seaborn) (3.10.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.6)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (26.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.3)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.27.1)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.2->seaborn) (2025.3)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)

```

```

# =====
# 3. Device & Hyperparameters
# =====

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

BATCH_SIZE = 32 # Compromise between originals (16 for ViT, 64 for ResNet)
NUM_EPOCHS = 30 # Balanced for comparison
NUM_CLASSES = 200

models_config = [
    {"name": "resnet18", "pretrained": True, "lr": 0.001},
    {"name": "vit", "pretrained": True, "lr": 3e-5}
]

```

Using device: cuda

```

# =====
# 4. Load Dataset from Hugging Face
# =====
print("Loading dataset...")
hf_dataset = load_dataset("bentrevett/caltech-ucsd-birds-200-2011")

class_names = hf_dataset['train'].features['label'].names # Get class names for reporting

class CUBDataset(torch.utils.data.Dataset):
    def __init__(self, hf_dataset, transform=None):
        self.hf_dataset = hf_dataset
        self.transform = transform

    def __len__(self):
        return len(self.hf_dataset)

    def __getitem__(self, idx):
        item = self.hf_dataset[idx]
        image = item['image'].convert('RGB')

```

```

        label = item['label']
        if self.transform:
            image = self.transform(image)
        return image, label

train_transform = transforms.Compose([
    transforms.Resize(256),
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(15),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

test_transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

full_train_ds = CUBDataset(hf_dataset['train'], train_transform)
test_ds = CUBDataset(hf_dataset['test'], test_transform)

train_size = int(0.85 * len(full_train_ds))
val_size = len(full_train_ds) - train_size
train_ds, val_ds = random_split(full_train_ds, [train_size, val_size])

train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE, shuffle=True, num_workers=2, pin_memory=True)
val_loader = DataLoader(val_ds, batch_size=BATCH_SIZE, shuffle=False, num_workers=2, pin_memory=True)
test_loader = DataLoader(test_ds, batch_size=BATCH_SIZE, shuffle=False, num_workers=2, pin_memory=True)

print(f"Train samples: {len(train_ds)}, Val samples: {len(val_ds)}, Test samples: {len(test_ds)}")
print(f"Num classes: {NUM_CLASSES}, Example classes: {class_names[:5]}...")

```

Loading dataset...
 Train samples: 5094, Val samples: 900, Test samples: 5794
 Num classes: 200, Example classes: ['001.Black_footed_Albatross', '002.Laysan_Albatross', '003.Sooty_Albatross', '004.Groove_billed_Ani']

```

# =====
# 5. Train and Evaluate Functions (Unified for both models)
# =====

def train_epoch(model, loader, criterion, optimizer, scaler, device, model_name):
    model.train()
    total_loss, correct, total = 0.0, 0, 0
    for images, labels in tqdm(loader, desc="Training"):
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        with torch.cuda.amp.autocast(enabled=True):
            if model_name == "vit":
                outputs = model(pixel_values=images).logits
            else:
                outputs = model(images)
            loss = criterion(outputs, labels)
            scaler.scale(loss).backward()
            scaler.step(optimizer)
            scaler.update()
            total_loss += loss.item() * images.size(0)
            pred = outputs.argmax(dim=1)
            correct += (pred == labels).sum().item()
            total += images.size(0)
    return total_loss / total, correct / total

def evaluate(model, loader, criterion, device, model_name, is_test=False):
    model.eval()
    total_loss, correct, total = 0.0, 0, 0
    all_preds, all_labels, all_logits = [], [], []
    with torch.no_grad():
        for images, labels in tqdm(loader, desc="Evaluating"):
            images, labels = images.to(device), labels.to(device)
            if model_name == "vit":
                outputs = model(pixel_values=images).logits
            else:
                outputs = model(images)
            loss = criterion(outputs, labels)
            total_loss += loss.item() * images.size(0)
            pred = outputs.argmax(dim=1)
            correct += (pred == labels).sum().item()
            total += images.size(0)
            all_preds.append(pred.cpu().numpy())
            all_labels.append(labels.cpu().numpy())
            all_logits.append(outputs.cpu().numpy())
    return total_loss / total, correct / total, all_preds, all_labels, all_logits

```

```

        all_labels.extend(labels.cpu().numpy())
        if is_test:
            all_logits.extend(outputs.cpu().numpy())    # For top-5
    acc = correct / total
    loss = total_loss / total
    top5_acc = top_k_accuracy_score(all_labels, all_logits, k=5) if is_test else None
    return loss, acc, top5_acc, all_preds, all_labels

```

```

# =====
# 6. Main Training Loop
# =====
all_history = {}
test_results = {}
for config in models_config:
    name = config["name"]
    print(f"\n== Training {name} ==")
    if name == "resnet18":
        model = models.resnet18(pretrained=config["pretrained"])
        model.fc = nn.Linear(model.fc.in_features, NUM_CLASSES)
        # Freeze all except last layer and classifier (as in original ResNet)
        for param in model.parameters():
            param.requires_grad = False
        for param in model.layer4.parameters():
            param.requires_grad = True
        for param in model.fc.parameters():
            param.requires_grad = True

    elif name == "vit":
        model = ViTForImageClassification.from_pretrained(
            "google/vit-base-patch16-224-in21k",
            num_labels=NUM_CLASSES,
            ignore_mismatched_sizes=True
        )
        # No freezing in original ViT, so train all with small LR

    model = model.to(device)
    criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
    optimizer = optim.AdamW(model.parameters(), lr=config["lr"], weight_decay=0.01)
    scaler = torch.cuda.amp.GradScaler(enabled=True)
    scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='max', factor=0.5, patience=2)

    history = {"train_loss": [], "train_acc": [], "val_loss": [], "val_acc": [], "train_time": []}
    best_val_acc = 0.0
    t_start = time.time()

    for epoch in range(NUM_EPOCHS):
        tr_loss, tr_acc = train_epoch(model, train_loader, criterion, optimizer, scaler, device, name)
        va_loss, va_acc, _, _, _ = evaluate(model, val_loader, criterion, device, name)
        scheduler.step(va_acc)
        history["train_loss"].append(tr_loss)
        history["train_acc"].append(tr_acc)
        history["val_loss"].append(va_loss)
        history["val_acc"].append(va_acc)
        epoch_time = time.time() - t_start
        history["train_time"].append(epoch_time)
        print(f"Epoch {epoch+1}/{NUM_EPOCHS} | Train Loss: {tr_loss:.4f} Acc: {tr_acc:.4f} | Val Loss: {va_loss:.4f}")

        if va_acc > best_val_acc:
            best_val_acc = va_acc
            torch.save(model.state_dict(), f"best_{name}.pth")

    total_train_time = time.time() - t_start
    print(f"Total training time for {name}: {total_train_time:.1f}s")

    # Load best model for test
    model.load_state_dict(torch.load(f"best_{name}.pth", map_location=device))
    te_loss, te_acc, te_top5_acc, te_preds, te_labels = evaluate(model, test_loader, criterion, device, name, is_test=True)
    report = classification_report(te_labels, te_preds, target_names=class_names, output_dict=True)
    macro_f1 = report['macro avg']['f1-score']

    test_results[name] = {
        "acc": te_acc,
        "top5_acc": te_top5_acc,
        "macro_f1": macro_f1,
        "train_time": total_train_time,
        "preds": te_preds,
        "labels": te_labels
    }
    all_history[name] = history

# Per-model error analysis (top-5 confused pairs)

```

```
errors = [(true, pred) for true, pred in zip(te_labels, te_preds) if true != pred]
if errors:
    most_common_errors = Counter(errors).most_common(5)
    print(f"Top 5 confused pairs for {name} (True → Predicted):")
    for (true_label, pred_label), count in most_common_errors:
        true_name = class_names[true_label]
        pred_name = class_names[pred_label]
        print(f"    {true_name} → {pred_name} : {count} times")
```



```
== Training resnet18 ==
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.
  warnings.warn(
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for
  warnings.warn(msg)
/tmp/ipython-input-3482151153.py:31: FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated. Please use `torch.amp.GradScaler
  scaler = torch.cuda.amp.GradScaler(enabled=True)
Training:  0% |                                     | 0/160 [00:00<?, ?it/s] /tmp/ipython-input-838357888.py:10: FutureWarning: `torch.cuda.amp.autocast(args...)` i
  with torch.cuda.amp.autocast(enabled=True):
Training: 100%|██████████| 160/160 [00:59<00:00,  2.67it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.16it/s]
Epoch 1/30 | Train Loss: 4.4131 Acc: 0.1227 | Val Loss: 3.6645 Acc: 0.2500
Training: 100%|██████████| 160/160 [00:32<00:00,  4.98it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.32it/s]
Epoch 2/30 | Train Loss: 3.3113 Acc: 0.3147 | Val Loss: 3.2139 Acc: 0.3178
Training: 100%|██████████| 160/160 [00:32<00:00,  4.93it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.32it/s]
Epoch 3/30 | Train Loss: 2.9849 Acc: 0.4009 | Val Loss: 3.1517 Acc: 0.3944
Training: 100%|██████████| 160/160 [00:35<00:00,  4.52it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.48it/s]
Epoch 4/30 | Train Loss: 2.7822 Acc: 0.4609 | Val Loss: 3.0364 Acc: 0.3944
Training: 100%|██████████| 160/160 [00:34<00:00,  4.69it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  5.18it/s]
Epoch 5/30 | Train Loss: 2.5847 Acc: 0.5261 | Val Loss: 2.8577 Acc: 0.4511
Training: 100%|██████████| 160/160 [00:35<00:00,  4.49it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  5.22it/s]
Epoch 6/30 | Train Loss: 2.4598 Acc: 0.5626 | Val Loss: 2.8177 Acc: 0.4744
Training: 100%|██████████| 160/160 [00:33<00:00,  4.77it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  5.25it/s]
Epoch 7/30 | Train Loss: 2.3739 Acc: 0.5885 | Val Loss: 2.7621 Acc: 0.4811
Training: 100%|██████████| 160/160 [00:33<00:00,  4.81it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  5.07it/s]
Epoch 8/30 | Train Loss: 2.2520 Acc: 0.6229 | Val Loss: 2.7941 Acc: 0.4767
Training: 100%|██████████| 160/160 [00:35<00:00,  4.48it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  5.06it/s]
Epoch 9/30 | Train Loss: 2.2141 Acc: 0.6378 | Val Loss: 2.7987 Acc: 0.4822
Training: 100%|██████████| 160/160 [00:33<00:00,  4.77it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.66it/s]
Epoch 10/30 | Train Loss: 2.1279 Acc: 0.6639 | Val Loss: 2.8092 Acc: 0.4744
Training: 100%|██████████| 160/160 [00:33<00:00,  4.85it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.35it/s]
Epoch 11/30 | Train Loss: 2.0650 Acc: 0.6959 | Val Loss: 2.7566 Acc: 0.4922
Training: 100%|██████████| 160/160 [00:32<00:00,  4.89it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.32it/s]
Epoch 12/30 | Train Loss: 2.0473 Acc: 0.6985 | Val Loss: 2.8087 Acc: 0.4767
Training: 100%|██████████| 160/160 [00:32<00:00,  4.88it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.30it/s]
Epoch 13/30 | Train Loss: 2.0067 Acc: 0.7087 | Val Loss: 2.7345 Acc: 0.5167
Training: 100%|██████████| 160/160 [00:32<00:00,  4.86it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.25it/s]
Epoch 14/30 | Train Loss: 1.9838 Acc: 0.7191 | Val Loss: 2.8173 Acc: 0.4822
Training: 100%|██████████| 160/160 [00:32<00:00,  4.90it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.36it/s]
Epoch 15/30 | Train Loss: 1.9302 Acc: 0.7346 | Val Loss: 2.7303 Acc: 0.5144
Training: 100%|██████████| 160/160 [00:32<00:00,  4.91it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.21it/s]
Epoch 16/30 | Train Loss: 1.8801 Acc: 0.7532 | Val Loss: 2.7800 Acc: 0.4989
Training: 100%|██████████| 160/160 [00:33<00:00,  4.81it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.37it/s]
Epoch 17/30 | Train Loss: 1.7388 Acc: 0.8074 | Val Loss: 2.5005 Acc: 0.5678
Training: 100%|██████████| 160/160 [00:32<00:00,  4.94it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.35it/s]
Epoch 18/30 | Train Loss: 1.6760 Acc: 0.8204 | Val Loss: 2.5576 Acc: 0.5467
Training: 100%|██████████| 160/160 [00:32<00:00,  4.94it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.28it/s]
Epoch 19/30 | Train Loss: 1.6369 Acc: 0.8316 | Val Loss: 2.6019 Acc: 0.5567
Training: 100%|██████████| 160/160 [00:32<00:00,  4.96it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.33it/s]
Epoch 20/30 | Train Loss: 1.6554 Acc: 0.8251 | Val Loss: 2.5767 Acc: 0.5544
Training: 100%|██████████| 160/160 [00:35<00:00,  4.53it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  4.98it/s]
Epoch 21/30 | Train Loss: 1.5740 Acc: 0.8543 | Val Loss: 2.5261 Acc: 0.5878
Training: 100%|██████████| 160/160 [00:35<00:00,  4.50it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  5.01it/s]
Epoch 22/30 | Train Loss: 1.5183 Acc: 0.8679 | Val Loss: 2.5261 Acc: 0.5744
Training: 100%|██████████| 160/160 [00:34<00:00,  4.67it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  5.17it/s]
Epoch 23/30 | Train Loss: 1.5095 Acc: 0.8687 | Val Loss: 2.5133 Acc: 0.5700
Training: 100%|██████████| 160/160 [00:34<00:00,  4.65it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  5.11it/s]
Epoch 24/30 | Train Loss: 1.5183 Acc: 0.8644 | Val Loss: 2.5986 Acc: 0.5667
Training: 100%|██████████| 160/160 [00:33<00:00,  4.71it/s]
Evaluating: 100%|██████████| 29/29 [00:05<00:00,  4.90it/s]
Epoch 25/30 | Train Loss: 1.5080 Acc: 0.8724 | Val Loss: 2.5243 Acc: 0.5800
```

```

Training: 100%|██████████| 160/160 [00:34<00:00,  4.66it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.19it/s]
Epoch 26/30 | Train Loss: 1.4747 Acc: 0.8820 | Val Loss: 2.5739 Acc: 0.5578
Training: 100%|██████████| 160/160 [00:33<00:00,  4.81it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.28it/s]
Epoch 27/30 | Train Loss: 1.4515 Acc: 0.8816 | Val Loss: 2.5614 Acc: 0.5733
Training: 100%|██████████| 160/160 [00:33<00:00,  4.84it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.18it/s]
Epoch 28/30 | Train Loss: 1.4677 Acc: 0.8783 | Val Loss: 2.5488 Acc: 0.5578
Training: 100%|██████████| 160/160 [00:33<00:00,  4.84it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.28it/s]
Epoch 29/30 | Train Loss: 1.4435 Acc: 0.8850 | Val Loss: 2.6222 Acc: 0.5511
Training: 100%|██████████| 160/160 [00:33<00:00,  4.78it/s]
Evaluating: 100%|██████████| 29/29 [00:06<00:00,  4.53it/s]
Epoch 30/30 | Train Loss: 1.4342 Acc: 0.8912 | Val Loss: 2.5862 Acc: 0.5611
Total training time for resnet18: 1225.9s
Evaluating: 100%|██████████| 182/182 [00:30<00:00,  6.01it/s]
Top 5 confused pairs for resnet18 (True → Predicted):
 025.Pelagic_Cormorant → 023.Brandt_Cormorant : 14 times
 051.Horned_Grebe → 050.Eared_Grebe : 12 times
 107.Common_Raven → 029.American_Crow : 11 times
 185.Bohemian_Waxwing → 186.Cedar_Waxwing : 10 times
 102.Western_Wood_Pewee → 040.Olive_sided_Flycatcher : 9 times

```

==== Training vit ===

Loading weights: 100% 198/198 [00:00<00:00, 399.68it/s, Materializing param=vit.layernorm.weight]

ViTForImageClassification LOAD REPORT from: google/vit-base-patch16-224-in21k

Key	Status
pooler.dense.weight	UNEXPECTED
pooler.dense.bias	UNEXPECTED
classifier.bias	MISSING
classifier.weight	MISSING

Notes:

- UNEXPECTED : can be ignored when loading from different task/architecture; not ok if you expect identical arch.
 - MISSING : those params were newly initialized because missing from the checkpoint. Consider training on your downstream task.
- /tmp/ipython-input-3482151153.py:31: FutureWarning: `torch.cuda.amp.GradScaler(args...)` is deprecated. Please use `torch.amp.GradScaler`
- scaler = torch.cuda.amp.GradScaler(enabled=True)
- Training: 0%| 0/160 [00:00<?, ?it/s]/tmp/ipython-input-838357888.py:10: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast` with torch.cuda.amp.autocast(enabled=True):
- Training: 100%|██████████| 160/160 [00:48<00:00, 3.33it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.74it/s]
Epoch 1/30 | Train Loss: 5.1987 Acc: 0.0870 | Val Loss: 5.0642 Acc: 0.2622
Training: 100%|██████████| 160/160 [00:47<00:00, 3.34it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.77it/s]
Epoch 2/30 | Train Loss: 4.8882 Acc: 0.4321 | Val Loss: 4.7495 Acc: 0.5022
Training: 100%|██████████| 160/160 [00:47<00:00, 3.38it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.80it/s]
Epoch 3/30 | Train Loss: 4.5550 Acc: 0.5850 | Val Loss: 4.4486 Acc: 0.5811
Training: 100%|██████████| 160/160 [00:46<00:00, 3.41it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.75it/s]
Epoch 4/30 | Train Loss: 4.2411 Acc: 0.6466 | Val Loss: 4.1845 Acc: 0.5711
Training: 100%|██████████| 160/160 [00:47<00:00, 3.38it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.82it/s]
Epoch 5/30 | Train Loss: 3.9398 Acc: 0.6920 | Val Loss: 3.8888 Acc: 0.6311
Training: 100%|██████████| 160/160 [00:47<00:00, 3.40it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.73it/s]
Epoch 6/30 | Train Loss: 3.6514 Acc: 0.7189 | Val Loss: 3.6511 Acc: 0.6344
Training: 100%|██████████| 160/160 [00:48<00:00, 3.33it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.78it/s]
Epoch 7/30 | Train Loss: 3.3669 Acc: 0.7313 | Val Loss: 3.3957 Acc: 0.6467
Training: 100%|██████████| 160/160 [00:47<00:00, 3.37it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.75it/s]
Epoch 8/30 | Train Loss: 3.0898 Acc: 0.7585 | Val Loss: 3.1805 Acc: 0.6378
Training: 100%|██████████| 160/160 [00:47<00:00, 3.35it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.79it/s]
Epoch 9/30 | Train Loss: 2.8423 Acc: 0.7774 | Val Loss: 2.9149 Acc: 0.6656
Training: 100%|██████████| 160/160 [00:47<00:00, 3.37it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.77it/s]
Epoch 10/30 | Train Loss: 2.6150 Acc: 0.7880 | Val Loss: 2.7135 Acc: 0.6867
Training: 100%|██████████| 160/160 [00:47<00:00, 3.35it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.78it/s]
Epoch 11/30 | Train Loss: 2.3871 Acc: 0.8078 | Val Loss: 2.5997 Acc: 0.6800
Training: 100%|██████████| 160/160 [00:47<00:00, 3.40it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.80it/s]
Epoch 12/30 | Train Loss: 2.2042 Acc: 0.8157 | Val Loss: 2.3435 Acc: 0.7267
Training: 100%|██████████| 160/160 [00:47<00:00, 3.39it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.80it/s]
Epoch 13/30 | Train Loss: 2.0458 Acc: 0.8245 | Val Loss: 2.3134 Acc: 0.7100
Training: 100%|██████████| 160/160 [00:47<00:00, 3.38it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.80it/s]
Epoch 14/30 | Train Loss: 1.8975 Acc: 0.8430 | Val Loss: 2.1997 Acc: 0.7167
Training: 100%|██████████| 160/160 [00:47<00:00, 3.40it/s]
Evaluating: 100%|██████████| 29/29 [00:10<00:00, 2.77it/s]

```

# =====
# 7. Compare Models
# =====
print("\n== Model Comparison ==")
comparison_data = {
    "Model": list(test_results.keys()),
    "Test Accuracy": [f"{test_results[m]['acc']:.4f}" for m in test_results],
    "Test Top-5 Accuracy": [f"{test_results[m]['top5_acc']:.4f}" for m in test_results],
    "Macro F1-Score": [f"{test_results[m]['macro_f1']:.4f}" for m in test_results],
    "Total Train Time (s)": [f"{test_results[m]['train_time']:.1f}" for m in test_results]
}
print(pd.DataFrame(comparison_data).to_markdown(index=False)) # Requires import pandas as pd if needed

# Plot combined curves
plt.figure(figsize=(12, 5))
for name in all_history:
    plt.subplot(1, 2, 1)
    plt.plot(all_history[name]["train_loss"], label=f"{name} train_loss")
    plt.plot(all_history[name]["val_loss"], label=f"{name} val_loss")
    plt.title("Loss Curves")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.legend()

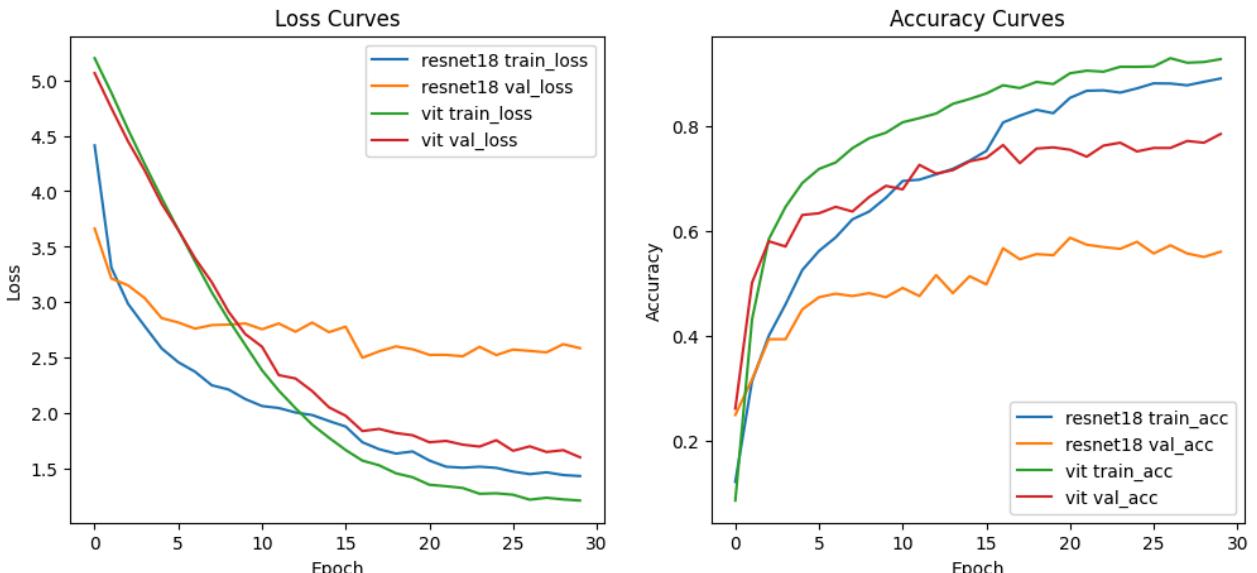
    plt.subplot(1, 2, 2)
    plt.plot(all_history[name]["train_acc"], label=f"{name} train_acc")
    plt.plot(all_history[name]["val_acc"], label=f"{name} val_acc")
    plt.title("Accuracy Curves")
    plt.xlabel("Epoch")
    plt.ylabel("Accuracy")
    plt.legend()
plt.show()

```

```

Evaluating: 100% |██████████| 29/29 [00:10<00:00, 2.74it/s]
Epoch 27/30 | Train Loss: 1.2215 Acc: 0.9299 | Val Loss: 1.7018 Acc: 0.7589
Training: 100% |██████████| 160/160 [00:48<00:00, 3.27it/s]
Evaluating: 100% |██████████| 29/29 [00:00<00:00, 0.00it/s] Macro F1-Score | Total Train Time (s) |
Epoch 28/30 | Train Loss: 1.2386 Acc: 0.9213 | Val Loss: 1.6510 Acc: 0.7722
Training: 100% |██████████| 160/160 [00:48<00:00, 3.30it/s] Macro F1-Score | Total Train Time (s) |
Evaluating: 100% |██████████| 29/29 [00:00<00:00, 0.00it/s] Macro F1-Score | Total Train Time (s) |

```



```

# =====
# 8. Plot Confusion Matrices for Both Models
# =====
print("\n== Plotting Confusion Matrices ==")

for name in test_results:
    y_true = test_results[name]['labels'] # list or np.array of true labels
    y_pred = test_results[name]['preds'] # list or np.array of predicted labels

    # Compute Confusion Matrice (200x200)
    cm = confusion_matrix(y_true, y_pred)

    # Normalized (row-normalized to 1, easier to see error distribution)
    cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
    cm_normalized = np.nan_to_num(cm_normalized)

```